

Machine Learning I
80-629A


Apprentissage Automatique I
80-629

Sequential Decision Making II
— Week #12

Introduction to Reinforcement Learning

Brief recap

- Markov Decision Processes (MDP)
- Offer a framework for sequential decision making
 $\langle \mathbf{A}, \mathbf{S}, \mathbf{P}, \mathbf{R}, \gamma \rangle$
- Goal: find the optimal policy
- Dynamic programming and several algorithms (e.g., VI,PI)



From MDPs to RL

- In MDPs we assume that we know

From MDPs to RL

- In MDPs we assume that we know
 1. Transition probabilities: $P(s' | s, a)$

From MDPs to RL

- In MDPs we assume that we know
 1. Transition probabilities: $P(s' | s, a)$
 2. Reward function: $R(s)$

From MDPs to RL

- In MDPs we assume that we know
 1. Transition probabilities: $P(s' | s, a)$
 2. Reward function: $R(s)$
- RL is more general

From MDPs to RL

- In MDPs we assume that we know
 1. Transition probabilities: $P(s' | s, a)$
 2. Reward function: $R(s)$
- RL is more general
 - In RL both are typically unknown

From MDPs to RL

- In MDPs we assume that we know
 1. Transition probabilities: $P(s' | s, a)$
 2. Reward function: $R(s)$
- RL is more general
 - In RL both are typically unknown
 - RL agents navigate the world to gather this information

Experience

A. Supervised Learning:

- Given fixed dataset
- Goal: maximize objective on test set (population)

B. Reinforcement Learning

- Collect data as agent interacts with the world
- Goal: maximize sum of rewards

RL applications

- **Key:** decision making over time, uncertain environments

RL applications

- **Key:** decision making over time, uncertain environments
- **Robot navigation:** Self-driving cars, helicopter control

RL applications

- **Key:** decision making over time, uncertain environments
- **Robot navigation:** Self-driving cars, helicopter control
- **Interactive systems:** recommender systems, chatbots

RL applications


- **Key:** decision making over time, uncertain environments
- **Robot navigation:** Self-driving cars, helicopter control
- **Interactive systems:** recommender systems, chatbots
- **Game playing:** Backgammon, go

RL applications

- **Key:** decision making over time, uncertain environments
- **Robot navigation:** Self-driving cars, helicopter control
- **Interactive systems:** recommender systems, chatbots
- **Game playing:** Backgammon, go
- **Healthcare:** monitoring systems

Reinforcement learning and recommender systems

- Most users have multiple interactions with the system over time
- Making recommendations over time can be advantageous (e.g., you could better explore one's preferences)
- States: Some representation of user preferences (e.g., previous items they consumed)
- Actions: what to recommend (item 1, item 2, item 3, ...)
- Reward:
 - + user consumes the recommendation
 - - user does not consume the recommendation



Algorithms for Reinforcement Learning

Algorithm

Algorithm

- **Input: an environment**
 - **actions, states, discount factor**
 - **starting state, method for obtaining next state**

Algorithm

- **Input: an environment**
 - **actions, states, discount factor**
 - **starting state, method for obtaining next state**
- **Output: an optimal policy**

Algorithm

- Input: an environment
 - actions, states, discount factor
 - starting state, method for obtaining next state
- Output: an optimal policy
- In practice: need a simulator or a real environment for your agent to interact

Algorithms for RL

- Two main classes of approach

Algorithms for RL

- Two main classes of approach

1. Model-based

- Learns a model of the transition and uses it to optimize a policy given the model

$$P(s' | s, a)$$

Algorithms for RL

- Two main classes of approach

1. Model-based

- Learns a model of the transition and uses it to optimize a policy given the model

$$P(s' | s, a)$$

2. Model-free

- Learns an optimal policy without explicitly learning transitions

$$\pi$$

Monte Carlo Methods

Monte Carlo Methods

- Model-free

Monte Carlo Methods

- Model-free
- Assume the environment is episodic
 - Think of playing a card game (like poker). An episode is a hand.
- Updates the policy after each episode

Monte Carlo Methods

- Model-free
- Assume the environment is episodic
 - Think of playing a card game (like poker). An episode is a hand.
- Updates the policy after each episode
- Intuition
 - Experience many episodes
 - Play many hands (of poker)
 - Average the rewards received at each state
 - What is the proportion of wins given your current cards

Prediction vs. control

1. Prediction: evaluate a given policy
2. Control: Learn a policy
 - Sometimes also called
 - passive (prediction)
 - active (control)

First-visit Monte Carlo

- Given a fixed policy (prediction)
- Calculate the value function $V(s)$ for each state

First-visit MC prediction, for estimating $V \approx v_\pi$

Initialize:

$\pi \leftarrow$ policy to be evaluated
 $V \leftarrow$ an arbitrary state-value function
 $Returns(s) \leftarrow$ an empty list, for all $s \in \mathcal{S}$

Repeat forever:

Generate an episode using π
For each state s appearing in the episode:
 $G \leftarrow$ the return that follows the first occurrence of s
 Append G to $Returns(s)$
 $V(s) \leftarrow \text{average}(Returns(s))$

[Sutton & Barto,
RL Book, Ch 5]

- Converges to $V_\pi(s)$ as the number of visits to each state goes to infinity

First-visit Monte Carlo

- Given a fixed policy (prediction)
- Calculate the value function $V(s)$ for each state

First-visit MC prediction, for estimating $V \approx v_\pi$

Initialize:

$\pi \leftarrow$ policy to be evaluated
 $V \leftarrow$ an arbitrary state-value function
 $Returns(s) \leftarrow$ an empty list, for all $s \in \mathcal{S}$

Repeat forever:

Generate an episode using π

For each state s appearing in the episode:

$G \leftarrow$ the return that follows the first occurrence of s

Append G to $Returns(s)$


$V(s) \leftarrow \text{average}(Returns(s))$

[Sutton & Barto,
RL Book, Ch 5]

- Converges to $V_\pi(s)$ as the number of visits to each state goes to infinity

$$V(s_t) = \max_{a_t} \left\{ R(s_t) + \gamma \sum_{s_{t+1}} P(s_{t+1} | s_t, a_t) V(s_{t+1}) \right\}$$

Example: grid world



- Start state is top-left (start of episode)
- Bottom right is absorbing (end of episode)
- Policy π is given (gray arrows)

Episode: (1, \rightarrow)

First-visit MC prediction, for estimating $V \approx v_\pi$

Initialize:

$\pi \leftarrow$ policy to be evaluated

$V \leftarrow$ an arbitrary state-value function

$Returns(s) \leftarrow$ an empty list, for all $s \in \mathcal{S}$

Repeat forever:

Generate an episode using π


For each state s appearing in the episode:

$G \leftarrow$ the return that follows the first occurrence of s

Append G to $Returns(s)$

$V(s) \leftarrow \text{average}(Returns(s))$

Example: grid world



- Start state is top-left (start of episode)
- Bottom right is absorbing (end of episode)
- Policy π is given (gray arrows)

Episode: (1, \rightarrow)

First-visit MC prediction, for estimating $V \approx v_\pi$

Initialize:

$\pi \leftarrow$ policy to be evaluated

$V \leftarrow$ an arbitrary state-value function

$Returns(s) \leftarrow$ an empty list, for all $s \in \mathcal{S}$

Repeat forever:

Generate an episode using π


For each state s appearing in the episode:

$G \leftarrow$ the return that follows the first occurrence of s

Append G to $Returns(s)$

$V(s) \leftarrow \text{average}(Returns(s))$

Example: grid world



- Start state is top-left (start of episode)
- Bottom right is absorbing (end of episode)
- Policy π is given (gray arrows)

Episode: (1, \rightarrow)

First-visit MC prediction, for estimating $V \approx v_\pi$

Initialize:

$\pi \leftarrow$ policy to be evaluated

$V \leftarrow$ an arbitrary state-value function

$Returns(s) \leftarrow$ an empty list, for all $s \in \mathcal{S}$

Repeat forever:

Generate an episode using π


For each state s appearing in the episode:

$G \leftarrow$ the return that follows the first occurrence of s

Append G to $Returns(s)$

$V(s) \leftarrow \text{average}(Returns(s))$

Example: grid world



- Start state is top-left (start of episode)
- Bottom right is absorbing (end of episode)
- Policy π is given (gray arrows)

Episode: (1, \rightarrow) \rightarrow (2, \rightarrow)

First-visit MC prediction, for estimating $V \approx v_\pi$

Initialize:

$\pi \leftarrow$ policy to be evaluated

$V \leftarrow$ an arbitrary state-value function

$Returns(s) \leftarrow$ an empty list, for all $s \in \mathcal{S}$

Repeat forever:

Generate an episode using π


For each state s appearing in the episode:

$G \leftarrow$ the return that follows the first occurrence of s

Append G to $Returns(s)$

$V(s) \leftarrow \text{average}(Returns(s))$

Example: grid world



- Start state is top-left (start of episode)
- Bottom right is absorbing (end of episode)
- Policy π is given (gray arrows)

Episode: (1, \rightarrow) \rightarrow (2, \rightarrow)

First-visit MC prediction, for estimating $V \approx v_\pi$

Initialize:

$\pi \leftarrow$ policy to be evaluated

$V \leftarrow$ an arbitrary state-value function

$Returns(s) \leftarrow$ an empty list, for all $s \in \mathcal{S}$

Repeat forever:

Generate an episode using π


For each state s appearing in the episode:

$G \leftarrow$ the return that follows the first occurrence of s

Append G to $Returns(s)$

$V(s) \leftarrow \text{average}(Returns(s))$

Example: grid world



- Start state is top-left (start of episode)
- Bottom right is absorbing (end of episode)
- Policy π is given (gray arrows)

Episode: (1, \rightarrow) \rightarrow (2, \rightarrow) \rightarrow (3, \downarrow)

First-visit MC prediction, for estimating $V \approx v_\pi$

Initialize:

$\pi \leftarrow$ policy to be evaluated

$V \leftarrow$ an arbitrary state-value function

$Returns(s) \leftarrow$ an empty list, for all $s \in \mathcal{S}$

Repeat forever:

Generate an episode using π


For each state s appearing in the episode:

$G \leftarrow$ the return that follows the first occurrence of s

Append G to $Returns(s)$

$V(s) \leftarrow \text{average}(Returns(s))$

Example: grid world



- Start state is top-left (start of episode)
- Bottom right is absorbing (end of episode)
- Policy π is given (gray arrows)

Episode: (1, \rightarrow) \rightarrow (2, \rightarrow) \rightarrow (3, \downarrow)

First-visit MC prediction, for estimating $V \approx v_\pi$

Initialize:

$\pi \leftarrow$ policy to be evaluated

$V \leftarrow$ an arbitrary state-value function

$Returns(s) \leftarrow$ an empty list, for all $s \in \mathcal{S}$

Repeat forever:

Generate an episode using π


For each state s appearing in the episode:

$G \leftarrow$ the return that follows the first occurrence of s

Append G to $Returns(s)$

$V(s) \leftarrow \text{average}(Returns(s))$

Example: grid world



- Start state is top-left (start of episode)
- Bottom right is absorbing (end of episode)
- Policy π is given (gray arrows)

Episode: $(1, \rightarrow) \rightarrow (2, \rightarrow) \rightarrow (3, \downarrow) \rightarrow (7, \downarrow)$

First-visit MC prediction, for estimating $V \approx v_\pi$

Initialize:

$\pi \leftarrow$ policy to be evaluated

$V \leftarrow$ an arbitrary state-value function

$Returns(s) \leftarrow$ an empty list, for all $s \in \mathcal{S}$

Repeat forever:

Generate an episode using π


For each state s appearing in the episode:

$G \leftarrow$ the return that follows the first occurrence of s

Append G to $Returns(s)$

$V(s) \leftarrow \text{average}(Returns(s))$

Example: grid world



- Start state is top-left (start of episode)
- Bottom right is absorbing (end of episode)
- Policy π is given (gray arrows)

Episode: $(1, \rightarrow) \rightarrow (2, \rightarrow) \rightarrow (3, \downarrow) \rightarrow (7, \downarrow)$

First-visit MC prediction, for estimating $V \approx v_\pi$

Initialize:

$\pi \leftarrow$ policy to be evaluated

$V \leftarrow$ an arbitrary state-value function

$Returns(s) \leftarrow$ an empty list, for all $s \in \mathcal{S}$

Repeat forever:

Generate an episode using π


For each state s appearing in the episode:

$G \leftarrow$ the return that follows the first occurrence of s

Append G to $Returns(s)$

$V(s) \leftarrow \text{average}(Returns(s))$

Example: grid world



- Start state is top-left (start of episode)
- Bottom right is absorbing (end of episode)
- Policy π is given (gray arrows)

Episode: $(1, \rightarrow) \rightarrow (2, \rightarrow) \rightarrow (3, \downarrow) \rightarrow (7, \downarrow) \rightarrow (6, \rightarrow)$

First-visit MC prediction, for estimating $V \approx v_\pi$

Initialize:

$\pi \leftarrow$ policy to be evaluated

$V \leftarrow$ an arbitrary state-value function

$Returns(s) \leftarrow$ an empty list, for all $s \in \mathcal{S}$

Repeat forever:

Generate an episode using π


For each state s appearing in the episode:

$G \leftarrow$ the return that follows the first occurrence of s

Append G to $Returns(s)$

$V(s) \leftarrow \text{average}(Returns(s))$

Example: grid world



- Start state is top-left (start of episode)
- Bottom right is absorbing (end of episode)
- Policy π is given (gray arrows)

Episode: $(1, \rightarrow) \rightarrow (2, \rightarrow) \rightarrow (3, \downarrow) \rightarrow (7, \downarrow) \rightarrow (6, \rightarrow)$

First-visit MC prediction, for estimating $V \approx v_\pi$

Initialize:

$\pi \leftarrow$ policy to be evaluated

$V \leftarrow$ an arbitrary state-value function

$Returns(s) \leftarrow$ an empty list, for all $s \in \mathcal{S}$

Repeat forever:

Generate an episode using π


For each state s appearing in the episode:

$G \leftarrow$ the return that follows the first occurrence of s

Append G to $Returns(s)$

$V(s) \leftarrow \text{average}(Returns(s))$

Example: grid world



- Start state is top-left (start of episode)
- Bottom right is absorbing (end of episode)
- Policy π is given (gray arrows)

Episode: (1, \rightarrow) \rightarrow (2, \rightarrow) \rightarrow (3, \downarrow) \rightarrow (7, \downarrow) \rightarrow (6, \rightarrow) \rightarrow (7, \downarrow)

First-visit MC prediction, for estimating $V \approx v_\pi$

Initialize:

$\pi \leftarrow$ policy to be evaluated

$V \leftarrow$ an arbitrary state-value function

$Returns(s) \leftarrow$ an empty list, for all $s \in \mathcal{S}$

Repeat forever:

Generate an episode using π


For each state s appearing in the episode:

$G \leftarrow$ the return that follows the first occurrence of s

Append G to $Returns(s)$

$V(s) \leftarrow \text{average}(Returns(s))$

Example: grid world



- Start state is top-left (start of episode)
- Bottom right is absorbing (end of episode)
- Policy π is given (gray arrows)

Episode: (1, \rightarrow) \rightarrow (2, \rightarrow) \rightarrow (3, \downarrow) \rightarrow (7, \downarrow) \rightarrow (6, \rightarrow) \rightarrow (7, \downarrow)

First-visit MC prediction, for estimating $V \approx v_\pi$

Initialize:

$\pi \leftarrow$ policy to be evaluated

$V \leftarrow$ an arbitrary state-value function

$Returns(s) \leftarrow$ an empty list, for all $s \in \mathcal{S}$

Repeat forever:

Generate an episode using π


For each state s appearing in the episode:

$G \leftarrow$ the return that follows the first occurrence of s

Append G to $Returns(s)$

$V(s) \leftarrow \text{average}(Returns(s))$

Example: grid world



- Start state is top-left (start of episode)
- Bottom right is absorbing (end of episode)
- Policy π is given (gray arrows)

Episode: (1, \rightarrow) \rightarrow (2, \rightarrow) \rightarrow (3, \downarrow) \rightarrow (7, \downarrow) \rightarrow (6, \rightarrow) \rightarrow (7, \downarrow) \rightarrow (10, \downarrow)

First-visit MC prediction, for estimating $V \approx v_\pi$

Initialize:

$\pi \leftarrow$ policy to be evaluated

$V \leftarrow$ an arbitrary state-value function

$Returns(s) \leftarrow$ an empty list, for all $s \in \mathcal{S}$

Repeat forever:

Generate an episode using π


For each state s appearing in the episode:

$G \leftarrow$ the return that follows the first occurrence of s

Append G to $Returns(s)$

$V(s) \leftarrow \text{average}(Returns(s))$

Example: grid world



- Start state is top-left (start of episode)
- Bottom right is absorbing (end of episode)
- Policy π is given (gray arrows)

Episode: (1, \rightarrow) \rightarrow (2, \rightarrow) \rightarrow (3, \downarrow) \rightarrow (7, \downarrow) \rightarrow (6, \rightarrow) \rightarrow (7, \downarrow) \rightarrow (10, \downarrow)

First-visit MC prediction, for estimating $V \approx v_\pi$

Initialize:

$\pi \leftarrow$ policy to be evaluated

$V \leftarrow$ an arbitrary state-value function

$Returns(s) \leftarrow$ an empty list, for all $s \in \mathcal{S}$

Repeat forever:

Generate an episode using π


For each state s appearing in the episode:

$G \leftarrow$ the return that follows the first occurrence of s

Append G to $Returns(s)$

$V(s) \leftarrow \text{average}(Returns(s))$

Example: grid world



- Start state is top-left (start of episode)
- Bottom right is absorbing (end of episode)
- Policy π is given (gray arrows)

Episode: (1, \rightarrow) \rightarrow (2, \rightarrow) \rightarrow (3, \downarrow) \rightarrow (7, \downarrow) \rightarrow (6, \rightarrow) \rightarrow (7, \downarrow) \rightarrow (10, \downarrow) \rightarrow (13, \downarrow)

First-visit MC prediction, for estimating $V \approx v_\pi$

Initialize:

$\pi \leftarrow$ policy to be evaluated

$V \leftarrow$ an arbitrary state-value function

$Returns(s) \leftarrow$ an empty list, for all $s \in \mathcal{S}$

Repeat forever:

Generate an episode using π


For each state s appearing in the episode:

$G \leftarrow$ the return that follows the first occurrence of s

Append G to $Returns(s)$

$V(s) \leftarrow \text{average}(Returns(s))$

Example: grid world



- Start state is top-left (start of episode)
- Bottom right is absorbing (end of episode)
- Policy π is given (gray arrows)

Episode: (1, \rightarrow) \rightarrow (2, \rightarrow) \rightarrow (3, \downarrow) \rightarrow (7, \downarrow) \rightarrow (6, \rightarrow) \rightarrow (7, \downarrow) \rightarrow (10, \downarrow) \rightarrow (13, \downarrow)

First-visit MC prediction, for estimating $V \approx v_\pi$

Initialize:

$\pi \leftarrow$ policy to be evaluated

$V \leftarrow$ an arbitrary state-value function

$Returns(s) \leftarrow$ an empty list, for all $s \in \mathcal{S}$

Repeat forever:

Generate an episode using π


For each state s appearing in the episode:

$G \leftarrow$ the return that follows the first occurrence of s

Append G to $Returns(s)$

$V(s) \leftarrow \text{average}(Returns(s))$

Example: grid world



- Start state is top-left (start of episode)
- Bottom right is absorbing (end of episode)
- Policy π is given (gray arrows)

Episode: (1, \rightarrow) \rightarrow (2, \rightarrow) \rightarrow (3, \downarrow) \rightarrow (7, \downarrow) \rightarrow (6, \rightarrow) \rightarrow (7, \downarrow) \rightarrow (10, \downarrow) \rightarrow (13, \downarrow) \rightarrow (17, \rightarrow)

First-visit MC prediction, for estimating $V \approx v_\pi$

Initialize:

$\pi \leftarrow$ policy to be evaluated

$V \leftarrow$ an arbitrary state-value function

$Returns(s) \leftarrow$ an empty list, for all $s \in \mathcal{S}$

Repeat forever:

Generate an episode using π


For each state s appearing in the episode:

$G \leftarrow$ the return that follows the first occurrence of s

Append G to $Returns(s)$

$V(s) \leftarrow \text{average}(Returns(s))$

Example: grid world



- Start state is top-left (start of episode)
- Bottom right is absorbing (end of episode)
- Policy π is given (gray arrows)

Episode: (1, \rightarrow) \rightarrow (2, \rightarrow) \rightarrow (3, \downarrow) \rightarrow (7, \downarrow) \rightarrow (6, \rightarrow) \rightarrow (7, \downarrow) \rightarrow (10, \downarrow) \rightarrow (13, \downarrow) \rightarrow (17, \rightarrow)

First-visit MC prediction, for estimating $V \approx v_\pi$

Initialize:

$\pi \leftarrow$ policy to be evaluated

$V \leftarrow$ an arbitrary state-value function

$Returns(s) \leftarrow$ an empty list, for all $s \in \mathcal{S}$

Repeat forever:

Generate an episode using π


For each state s appearing in the episode:

$G \leftarrow$ the return that follows the first occurrence of s

Append G to $Returns(s)$

$V(s) \leftarrow \text{average}(Returns(s))$

Example: grid world



- Start state is top-left (start of episode)
- Bottom right is absorbing (end of episode)
- Policy π is given (gray arrows)

First-visit MC prediction, for estimating $V \approx v_\pi$

Initialize:

$\pi \leftarrow$ policy to be evaluated

$V \leftarrow$ an arbitrary state-value function

$Returns(s) \leftarrow$ an empty list, for all $s \in \mathcal{S}$

Repeat forever:

Generate an episode using π

For each state s appearing in the episode:


$G \leftarrow$ the return that follows the first occurrence of s

Append G to $Returns(s)$

$V(s) \leftarrow \text{average}(Returns(s))$

Episode: $(1, \rightarrow) \rightarrow (2, \rightarrow) \rightarrow (3, \downarrow) \rightarrow (7, \downarrow) \rightarrow (6, \rightarrow) \rightarrow (7, \downarrow) \rightarrow (10, \downarrow) \rightarrow (13, \downarrow) \rightarrow (17, \rightarrow)$

Example: grid world



- Start state is top-left (start of episode)
- Bottom right is absorbing (end of episode)
- Policy π is given (gray arrows)

First-visit MC prediction, for estimating $V \approx v_\pi$

Initialize:

$\pi \leftarrow$ policy to be evaluated

$V \leftarrow$ an arbitrary state-value function

$Returns(s) \leftarrow$ an empty list, for all $s \in \mathcal{S}$

Repeat forever:

Generate an episode using π

For each state s appearing in the episode:


$G \leftarrow$ the return that follows the first occurrence of s

Append G to $Returns(s)$

$V(s) \leftarrow \text{average}(Returns(s))$

Episode: $(1, \rightarrow) \rightarrow (2, \rightarrow) \rightarrow (3, \downarrow) \rightarrow (7, \downarrow) \rightarrow (6, \rightarrow) \rightarrow (7, \downarrow) \rightarrow (10, \downarrow) \rightarrow (13, \downarrow) \rightarrow (17, \rightarrow)$

Example: grid world



- Start state is top-left (start of episode)
- Bottom right is absorbing (end of episode)
- Policy π is given (gray arrows)

First-visit MC prediction, for estimating $V \approx v_\pi$

Initialize:

$\pi \leftarrow$ policy to be evaluated

$V \leftarrow$ an arbitrary state-value function

$Returns(s) \leftarrow$ an empty list, for all $s \in \mathcal{S}$

Repeat forever:

Generate an episode using π

For each state s appearing in the episode:

$G \leftarrow$ the return that follows the first occurrence of s


Append G to $Returns(s)$

$V(s) \leftarrow \text{average}(Returns(s))$

Episode: $(1, \rightarrow) \rightarrow (2, \rightarrow) \rightarrow (3, \downarrow) \rightarrow (7, \downarrow) \rightarrow (6, \rightarrow) \rightarrow (7, \downarrow) \rightarrow (10, \downarrow) \rightarrow (13, \downarrow) \rightarrow (17, \rightarrow)$

For state 7:

Example: grid world



- Start state is top-left (start of episode)
- Bottom right is absorbing (end of episode)
- Policy π is given (gray arrows)

First-visit MC prediction, for estimating $V \approx v_\pi$

Initialize:

$\pi \leftarrow$ policy to be evaluated

$V \leftarrow$ an arbitrary state-value function

$Returns(s) \leftarrow$ an empty list, for all $s \in \mathcal{S}$

Repeat forever:

Generate an episode using π

For each state s appearing in the episode:

$G \leftarrow$ the return that follows the first occurrence of s


Append G to $Returns(s)$

$V(s) \leftarrow \text{average}(Returns(s))$

Episode: $(1, \rightarrow) \rightarrow (2, \rightarrow) \rightarrow (3, \downarrow) \rightarrow (7, \downarrow) \rightarrow (6, \rightarrow) \rightarrow (7, \downarrow) \rightarrow (10, \downarrow) \rightarrow (13, \downarrow) \rightarrow (17, \rightarrow)$

For state 7:
$$\begin{aligned} \text{return}(7) &= \gamma R(6) + \gamma^2 R(7) + \gamma^3 R(10) + \gamma^4 R(13) + \gamma^5 R(17) + \gamma^6 R(18) \\ &= \gamma^6 10 \end{aligned}$$

Example: grid world



- Start state is top-left (start of episode)
- Bottom right is absorbing (end of episode)
- Policy π is given (gray arrows)

First-visit MC prediction, for estimating $V \approx v_\pi$

Initialize:

$\pi \leftarrow$ policy to be evaluated

$V \leftarrow$ an arbitrary state-value function

$Returns(s) \leftarrow$ an empty list, for all $s \in \mathcal{S}$

Repeat forever:

Generate an episode using π

For each state s appearing in the episode:

$G \leftarrow$ the return that follows the first occurrence of s


Append G to $Returns(s)$

$V(s) \leftarrow \text{average}(Returns(s))$

Episode: $(1, \rightarrow) \rightarrow (2, \rightarrow) \rightarrow (3, \downarrow) \rightarrow (7, \downarrow) \rightarrow (6, \rightarrow) \rightarrow (7, \downarrow) \rightarrow (10, \downarrow) \rightarrow (13, \downarrow) \rightarrow (17, \rightarrow)$

For state 7:
$$\begin{aligned} \text{return}(7) &= \gamma R(6) + \gamma^2 R(7) + \gamma^3 R(10) + \gamma^4 R(13) + \gamma^5 R(17) + \gamma^6 R(18) \\ &= \gamma^6 10 \end{aligned}$$

Example: grid world



- Start state is top-left (start of episode)
- Bottom right is absorbing (end of episode)
- Policy π is given (gray arrows)

First-visit MC prediction, for estimating $V \approx v_\pi$

Initialize:

$\pi \leftarrow$ policy to be evaluated

$V \leftarrow$ an arbitrary state-value function

$Returns(s) \leftarrow$ an empty list, for all $s \in \mathcal{S}$

Repeat forever:

Generate an episode using π

For each state s appearing in the episode:

$G \leftarrow$ the return that follows the first occurrence of s

Append G to $Returns(s)$

$V(s) \leftarrow \text{average}(Returns(s))$

Episode: $(1, \rightarrow) \rightarrow (2, \rightarrow) \rightarrow (3, \downarrow) \rightarrow (7, \downarrow) \rightarrow (6, \rightarrow) \rightarrow (7, \downarrow) \rightarrow (10, \downarrow) \rightarrow (13, \downarrow) \rightarrow (17, \rightarrow)$

For state 7:
$$\begin{aligned} \text{return}(7) &= \gamma R(6) + \gamma^2 R(7) + \gamma^3 R(10) + \gamma^4 R(13) + \gamma^5 R(17) + \gamma^6 R(18) \\ &= \gamma^6 10 \end{aligned}$$

$$V(7) = \gamma^6 * 10$$

Summary

- Introduced terminology:
 - model based, model-free
- First algorithm for policy evaluation (First-visit MC)
- Compared to MDPs
 - We the agent now has to explore the world to evaluate its value function

Algorithms for RL Control

Q-value function for control

- We know about state-value functions $V(s)$

Q-value function for control

- We know about state-value functions $V(s)$
- If state transitions are known then they can be used to derive an optimal policy [recall value iteration]:

$$\boldsymbol{\pi}^*(\mathbf{s}) = \arg \max_{\mathbf{a}} \left\{ \mathbf{R}(\mathbf{s}) + \gamma \sum_{\mathbf{s}'} \mathbf{P}(\mathbf{s}' \mid \mathbf{s}, \mathbf{a}) \mathbf{V}^*(\mathbf{s}') \right\} \quad \forall \mathbf{s}$$

Q-value function for control

- We know about state-value functions $V(s)$
- If state transitions are known then they can be used to derive an optimal policy [recall value iteration]:

$$\boldsymbol{\pi}^*(\mathbf{s}) = \arg \max_{\mathbf{a}} \left\{ \mathbf{R}(\mathbf{s}) + \gamma \sum_{s'} \mathbf{P}(s' \mid \mathbf{s}, \mathbf{a}) \mathbf{V}^*(s') \right\} \quad \forall \mathbf{s}$$

- When state transitions are unknown what can we do?

Q-value function for control

- We know about state-value functions $V(s)$
- If state transitions are known then they can be used to derive an optimal policy [recall value iteration]:

$$\pi^*(s) = \arg \max_a \left\{ R(s) + \gamma \sum_{s'} P(s' | s, a) V^*(s') \right\} \quad \forall s$$

- When state transitions are unknown what can we do?
- $Q(s,a)$ the value function of a (state,action) pair

$$\pi^*(s) = \arg \max_a \{ Q^*(s, a) \} \quad \forall s$$

Monte Carlo ES (control)

[Sutton & Barto,
RL Book, Ch.5]

Monte Carlo ES (Exploring Starts), for estimating $\pi \approx \pi_*$

Initialize, for all $s \in \mathcal{S}$, $a \in \mathcal{A}(s)$:

$Q(s, a) \leftarrow$ arbitrary

$\pi(s) \leftarrow$ arbitrary

$Returns(s, a) \leftarrow$ empty list

Repeat forever:

Choose $S_0 \in \mathcal{S}$ and $A_0 \in \mathcal{A}(S_0)$ s.t. all pairs have probability > 0

Generate an episode starting from S_0, A_0 , following π

For each pair s, a appearing in the episode:

$G \leftarrow$ the return that follows the first occurrence of s, a

Append G to $Returns(s, a)$

$Q(s, a) \leftarrow \text{average}(Returns(s, a))$

For each s in the episode:

$\pi(s) \leftarrow \arg\max_a Q(s, a)$

First-visit MC prediction, for estimating $V \approx V_*$

Initialize:

$\pi \leftarrow$ policy to be evaluated

$V \leftarrow$ an arbitrary state-value function

$Returns(s) \leftarrow$ an empty list, for all $s \in \mathcal{S}$

Repeat forever:

Generate an episode using π

For each state s appearing in the episode:

$G \leftarrow$ the return that follows the first occurrence

Append G to $Returns(s)$

$V(s) \leftarrow \text{average}(Returns(s))$

Monte Carlo ES (control)

[Sutton & Barto,
RL Book, Ch.5]

Monte Carlo ES (Exploring Starts), for estimating $\pi \approx \pi_*$

Initialize, for all $s \in \mathcal{S}$, $a \in \mathcal{A}(s)$:

$Q(s, a) \leftarrow$ arbitrary

$\pi(s) \leftarrow$ arbitrary

$Returns(s, a) \leftarrow$ empty list

Repeat forever:

Choose $S_0 \in \mathcal{S}$ and $A_0 \in \mathcal{A}(S_0)$ s.t. all pairs have probability > 0

Generate an episode starting from S_0, A_0 , following π

For each pair s, a appearing in the episode:

$G \leftarrow$ the return that follows the first occurrence of s, a

Append G to $Returns(s, a)$

$Q(s, a) \leftarrow \text{average}(Returns(s, a))$

For each s in the episode:

$\pi(s) \leftarrow \arg\max_a Q(s, a)$

First-visit MC prediction, for estimating $V \approx V_*$

Initialize:

$\pi \leftarrow$ policy to be evaluated

$V \leftarrow$ an arbitrary state-value function

$Returns(s) \leftarrow$ an empty list, for all $s \in \mathcal{S}$

Repeat forever:

Generate an episode using π

For each state s appearing in the episode:

$G \leftarrow$ the return that follows the first occurrence

Append G to $Returns(s)$

$V(s) \leftarrow \text{average}(Returns(s))$

- Strong reasons to believe that it converges to the optimal policy
- “Exploring starts” requirement may be unrealistic

Learning without “exploring starts”

- “Exploring starts” insures that all states can be visited regardless of the policy
- (Specific policy may not visit all states)
- Unrealistic in real-world settings

Learning without “exploring starts”

- “Exploring starts” insures that all states can be visited regardless of the policy
- (Specific policy may not visit all states)
- Unrealistic in real-world settings
- Solution: inject some uncertainty in the policy

Monte Carlo without exploring starts (on policy)

On-policy first-visit MC control (for ε -soft policies), estimates $\pi \approx \pi_*$

Initialize, for all $s \in \mathcal{S}$, $a \in \mathcal{A}(s)$:

$Q(s, a) \leftarrow$ arbitrary

$Returns(s, a) \leftarrow$ empty list

$\pi(a|s) \leftarrow$ an arbitrary ε -soft policy

Repeat forever:

(a) Generate an episode using π

(b) For each pair s, a appearing in the episode:

$G \leftarrow$ the return that follows the first occurrence of s, a

Append G to $Returns(s, a)$

$Q(s, a) \leftarrow \text{average}(Returns(s, a))$

(c) For each s in the episode:

$A^* \leftarrow \arg \max_a Q(s, a)$

(with ties broken arbitrarily)

For all $a \in \mathcal{A}(s)$:

$$\pi(a|s) \leftarrow \begin{cases} 1 - \varepsilon + \varepsilon/|\mathcal{A}(s)| & \text{if } a = A^* \\ \varepsilon/|\mathcal{A}(s)| & \text{if } a \neq A^* \end{cases}$$

[Sutton & Barto,
RL Book, Ch.5]

Monte Carlo ES (Exploring Starts), for estimating $\pi \approx \pi_*$

Initialize, for all $s \in \mathcal{S}$, $a \in \mathcal{A}(s)$:

$Q(s, a) \leftarrow$ arbitrary

$\pi(s) \leftarrow$ arbitrary

$Returns(s, a) \leftarrow$ empty list

Repeat forever:

Choose $S_0 \in \mathcal{S}$ and $A_0 \in \mathcal{A}(S_0)$ s.t. all pairs have probability > 0

Generate an episode starting from S_0, A_0 , following π

For each pair s, a appearing in the episode:

$G \leftarrow$ the return that follows the first occurrence of s, a

Append G to $Returns(s, a)$

$Q(s, a) \leftarrow \text{average}(Returns(s, a))$

For each s in the episode:

$\pi(s) \leftarrow \arg \max_a Q(s, a)$

Monte Carlo without exploring starts (on policy)

On-policy first-visit MC control (for ϵ -soft policies), estimates $\pi \approx \pi_*$

Initialize, for all $s \in \mathcal{S}$, $a \in \mathcal{A}(s)$:

$Q(s, a) \leftarrow$ arbitrary

$Returns(s, a) \leftarrow$ empty list

$\pi(a|s) \leftarrow$ an arbitrary ϵ -soft policy

Repeat forever:

(a) Generate an episode using π

(b) For each pair s, a appearing in the episode:

$G \leftarrow$ the return that follows the first occurrence of s, a

Append G to $Returns(s, a)$

$Q(s, a) \leftarrow \text{average}(Returns(s, a))$

(c) For each s in the episode:

$A^* \leftarrow \arg \max_a Q(s, a)$

(with ties broken arbitrarily)

For all $a \in \mathcal{A}(s)$:

$$\pi(a|s) \leftarrow \begin{cases} 1 - \epsilon + \epsilon/|\mathcal{A}(s)| & \text{if } a = A^* \\ \epsilon/|\mathcal{A}(s)| & \text{if } a \neq A^* \end{cases}$$

[Sutton & Barto,
RL Book, Ch.5]

Monte Carlo ES (Exploring Starts), for estimating $\pi \approx \pi_*$

Initialize, for all $s \in \mathcal{S}$, $a \in \mathcal{A}(s)$:

$Q(s, a) \leftarrow$ arbitrary

$\pi(s) \leftarrow$ arbitrary

$Returns(s, a) \leftarrow$ empty list

Repeat forever:

Choose $S_0 \in \mathcal{S}$ and $A_0 \in \mathcal{A}(S_0)$ s.t. all pairs have probability > 0

Generate an episode starting from S_0, A_0 , following π

For each pair s, a appearing in the episode:

$G \leftarrow$ the return that follows the first occurrence of s, a

Append G to $Returns(s, a)$

$Q(s, a) \leftarrow \text{average}(Returns(s, a))$

For each s in the episode:

$\pi(s) \leftarrow \arg \max_a Q(s, a)$

Monte Carlo without exploring starts (on policy)

On-policy first-visit MC control (for ϵ -soft policies), estimates $\pi \approx \pi_*$

Initialize, for all $s \in \mathcal{S}$, $a \in \mathcal{A}(s)$:

$Q(s, a) \leftarrow$ arbitrary

$Returns(s, a) \leftarrow$ empty list

$\pi(a|s) \leftarrow$ an arbitrary ϵ -soft policy

Repeat forever:

(a) Generate an episode using π

(b) For each pair s, a appearing in the episode:

$G \leftarrow$ the return that follows the first occurrence of s, a

Append G to $Returns(s, a)$

$Q(s, a) \leftarrow \text{average}(Returns(s, a))$

(c) For each s in the episode:

$A^* \leftarrow \arg \max_a Q(s, a)$

(with ties broken arbitrarily)

For all $a \in \mathcal{A}(s)$:

$$\pi(a|s) \leftarrow \begin{cases} 1 - \epsilon + \epsilon/|\mathcal{A}(s)| & \text{if } a = A^* \\ \epsilon/|\mathcal{A}(s)| & \text{if } a \neq A^* \end{cases}$$

[Sutton & Barto,
RL Book, Ch.5]

Monte Carlo ES (Exploring Starts), for estimating $\pi \approx \pi_*$

Initialize, for all $s \in \mathcal{S}$, $a \in \mathcal{A}(s)$:

$Q(s, a) \leftarrow$ arbitrary

$\pi(s) \leftarrow$ arbitrary

$Returns(s, a) \leftarrow$ empty list

Repeat forever:

Choose $S_0 \in \mathcal{S}$ and $A_0 \in \mathcal{A}(S_0)$ s.t. all pairs have probability > 0

Generate an episode starting from S_0, A_0 , following π

For each pair s, a appearing in the episode:

$G \leftarrow$ the return that follows the first occurrence of s, a

Append G to $Returns(s, a)$

$Q(s, a) \leftarrow \text{average}(Returns(s, a))$

For each s in the episode:

$\pi(s) \leftarrow \arg \max_a Q(s, a)$

Monte Carlo without exploring starts (on policy)

On-policy first-visit MC control (for ϵ -soft policies), estimates $\pi \approx \pi_*$

Initialize, for all $s \in \mathcal{S}$, $a \in \mathcal{A}(s)$:

$Q(s, a) \leftarrow$ arbitrary

$Returns(s, a) \leftarrow$ empty list

$\pi(a|s) \leftarrow$ an arbitrary ϵ -soft policy

Repeat forever:

(a) Generate an episode using π

(b) For each pair s, a appearing in the episode:

$G \leftarrow$ the return that follows the first occurrence of s, a

Append G to $Returns(s, a)$

$Q(s, a) \leftarrow \text{average}(Returns(s, a))$

(c) For each s in the episode:

$A^* \leftarrow \arg \max_a Q(s, a)$

(with ties broken arbitrarily)

For all $a \in \mathcal{A}(s)$:

$\pi(a|s) \leftarrow \begin{cases} 1 - \epsilon + \epsilon/|\mathcal{A}(s)| & \text{if } a = A^* \\ \epsilon/|\mathcal{A}(s)| & \text{if } a \neq A^* \end{cases}$

[Sutton & Barto,
RL Book, Ch.5]

Monte Carlo ES (Exploring Starts), for estimating $\pi \approx \pi_*$

Initialize, for all $s \in \mathcal{S}$, $a \in \mathcal{A}(s)$:

$Q(s, a) \leftarrow$ arbitrary

$\pi(s) \leftarrow$ arbitrary

$Returns(s, a) \leftarrow$ empty list

Repeat forever:

Choose $S_0 \in \mathcal{S}$ and $A_0 \in \mathcal{A}(S_0)$ s.t. all pairs have probability > 0

Generate an episode starting from S_0, A_0 , following π

For each pair s, a appearing in the episode:

$G \leftarrow$ the return that follows the first occurrence of s, a

Append G to $Returns(s, a)$

$Q(s, a) \leftarrow \text{average}(Returns(s, a))$

For each s in the episode:

$\pi(s) \leftarrow \arg \max_a Q(s, a)$

- Policy value cannot decrease

$$V_{\pi'}(s) \geq V_{\pi}(s), \forall s \in \mathcal{S}$$

π : policy at current step

π' : policy at next step

Monte-Carlo methods summary

- Allow a policy to be learned through interactions
 - (Does not learn transitions)
- States are effectively treated as being independent
 - Focus on a subset of states (e.g., states for which playing optimally is of particular importance)
- Episodic (with or without exploring starts)

Temporal Difference (TD) Learning

Temporal Difference (TD) Learning

- One of the “central ideas of RL” [Sutton & Barto, RL book]

Temporal Difference (TD) Learning

- One of the “central ideas of RL” [Sutton & Barto, RL book]

- Monte Carlo methods

$$\mathbf{V}'(\mathbf{s}_t) = \mathbf{V}(\mathbf{s}_t) + \alpha [\mathbf{G}_t - \mathbf{V}(\mathbf{s}_t)]$$

Observed return :

$$\mathbf{G}_t = \sum_t^T \gamma^t \mathbf{R}(\mathbf{s}_t)$$

Step size

Temporal Difference (TD) Learning

- One of the “central ideas of RL” [Sutton & Barto, RL book]

- Monte Carlo methods

$$V'(s_t) = V(s_t) + \alpha [G_t - V(s_t)]$$

Observed return :

Step size

$$G_t = \sum_t^T \gamma^t R(s_t)$$

First-visit MC prediction, for estimation

Initialize:

$\pi \leftarrow$ policy to be evaluated

$V \leftarrow$ an arbitrary state-value function

$Returns(s) \leftarrow$ an empty list, for all $s \in \mathcal{S}$

Repeat forever:

Generate an episode using π

For each state s appearing in the episode

$G \leftarrow$ the return that follows the first

Append G to $Returns(s)$

$V(s) \leftarrow \text{average}(Returns(s))$

Temporal Difference (TD) Learning

- One of the “central ideas of RL” [Sutton & Barto, RL book]

- Monte Carlo methods

$$V'(s_t) = V(s_t) + \alpha [G_t - V(s_t)]$$

Observed return :

Step size

$$G_t = \sum_t^T \gamma^t R(s_t)$$

- TD(0)

- updates “instantly”

First-visit MC prediction, for estimation

Initialize:

$\pi \leftarrow$ policy to be evaluated

$V \leftarrow$ an arbitrary state-value function

$Returns(s) \leftarrow$ an empty list, for all $s \in S$

Repeat forever:

Generate an episode using π

For each state s appearing in the episode

$G \leftarrow$ the return that follows the first

Append G to $Returns(s)$

$V(s) \leftarrow \text{average}(Returns(s))$

Temporal Difference (TD) Learning

- One of the “central ideas of RL” [Sutton & Barto, RL book]

- Monte Carlo methods

$$V'(s_t) = V(s_t) + \alpha [G_t - V(s_t)]$$

Observed return : G_t

Step size : α

$$G_t = \sum_t^T \gamma^t R(s_t)$$

- TD(0)

- updates “instantly”

$$V'(s_t) = V(s_t) + \alpha [\underbrace{R(s_t) + \gamma V(s_{t+1})}_{\approx G_t} - V(s_t)]$$

First-visit MC prediction, for estimation

Initialize:

$\pi \leftarrow$ policy to be evaluated

$V \leftarrow$ an arbitrary state-value function

$Returns(s) \leftarrow$ an empty list, for all $s \in S$

Repeat forever:

Generate an episode using π

For each state s appearing in the episode

$G \leftarrow$ the return that follows the first

Append G to $Returns(s)$

$V(s) \leftarrow \text{average}(Returns(s))$

TD(0) for prediction

[Sutton & Barto,
RL Book, Ch.6]

Tabular TD(0) for estimating v_π

```
Input: the policy  $\pi$  to be evaluated
Initialize  $V(s)$  arbitrarily (e.g.,  $V(s) = 0$ , for all  $s \in \mathcal{S}^+$ )
Repeat (for each episode):
  Initialize  $S$ 
  Repeat (for each step of episode):
     $A \leftarrow$  action given by  $\pi$  for  $S$ 
    Take action  $A$ , observe  $R, S'$ 
     $V(S) \leftarrow V(S) + \alpha[R + \gamma V(S') - V(S)]$ 
     $S \leftarrow S'$ 
  until  $S$  is terminal
```


TD for control

Sarsa (on-policy TD control) for estimating $Q \approx q_*$

Algorithm parameters: step size $\alpha \in (0, 1]$, small $\varepsilon > 0$

Initialize $Q(s, a)$, for all $s \in \mathcal{S}^+, a \in \mathcal{A}(s)$, arbitrarily except that $Q(\text{terminal}, \cdot) = 0$

Loop for each episode:

 Initialize S

 Choose A from S using policy derived from Q (e.g., ε -greedy)

 Loop for each step of episode:

 Take action A , observe R, S'

 Choose A' from S' using policy derived from Q (e.g., ε -greedy)

$Q(S, A) \leftarrow Q(S, A) + \alpha[R + \gamma Q(S', A') - Q(S, A)]$

$S \leftarrow S'; A \leftarrow A';$

 until S is terminal

Tabular TD(0) for estimating v_π

Input: the policy π to be evaluated

Initialize $V(s)$ arbitrarily (e.g., $V(s) = 0$, for all $s \in \mathcal{S}^+$)

Repeat (for each episode):

 Initialize S

 Repeat (for each step of episode):

$A \leftarrow$ action given by π for S

 Take action A , observe R, S'

$V(S) \leftarrow V(S) + \alpha[R + \gamma V(S') - V(S)]$

$S \leftarrow S'$

 until S is terminal

Comparing TD and MC

- MC requires going through full episodes before updating the value function. Episodic.
- Converges to the optimal solution
- TD updates each $V(s)$ after each transition. Online.
- Converges to the optimal solution (some conditions on α)
- Empirically TD methods tend to converge faster


Extra material
(Some will be used for
this week's exercises)

Example: Black Jack

- Episode: one hand
- States: Sum of player's cards, dealer's card, usable ace
- Actions: {Stay, Hit}
- Rewards: {Win +1, Tie 0, Loose -1}
- A few other assumptions: infinite deck

- Evaluates the policy that hits except when the sum of the cards is 20 or 21

No
usable
ace




[Figure 5.1, Sutton & Barto]

- Evaluates the policy that hits except when the sum of the cards is 20 or 21

Usable
ace

No
usable
ace




[Figure 5.1, Sutton & Barto]

- Evaluates the policy that hits except when the sum of the cards is 20 or 21

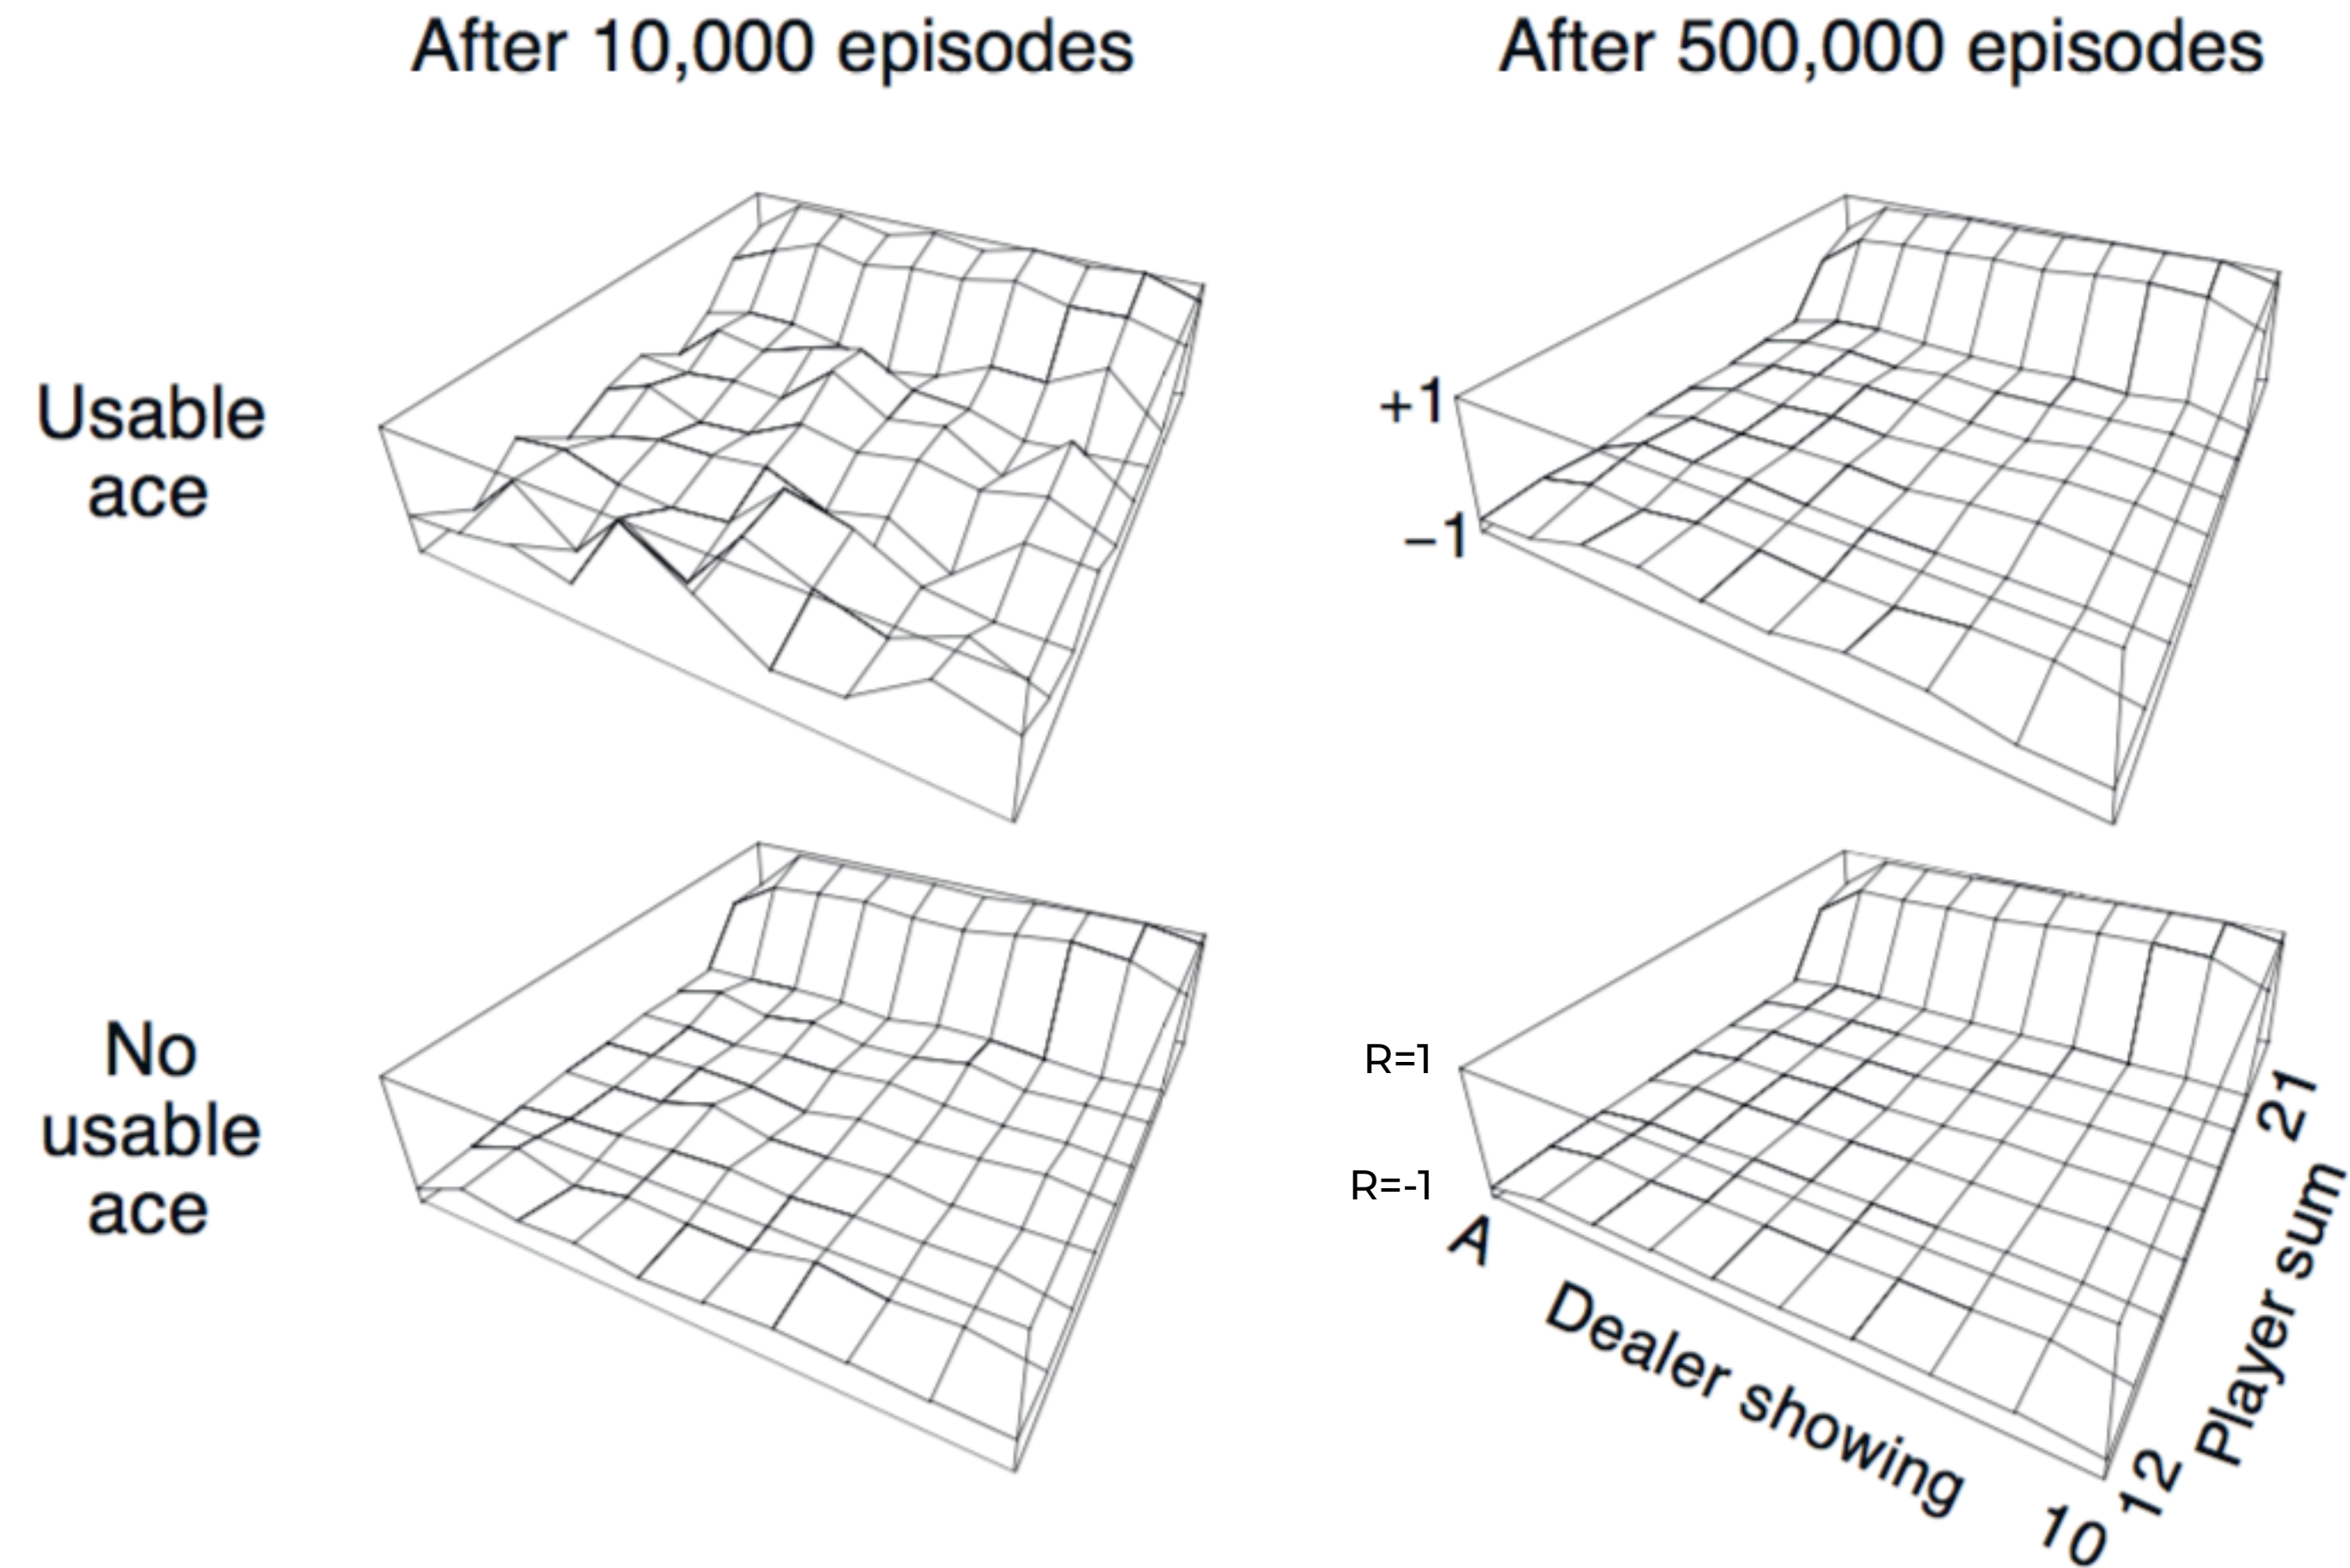
Usable
ace

No
usable
ace



[Figure 5.1, Sutton & Barto]

- Evaluates the policy that hits except when the sum of the cards is 20 or 21



[Figure 5.1, Sutton & Barto]

Practical difficulties

- Compared to supervised learning setting up an RL problem is often harder
 - Need an environment (or at least a simulator)
- Rewards
 - In some domains it's clear (e.g., in games)
 - In others it's much more subtle (e.g., you want to please a human)

Acknowledgements

- The algorithms are from “Reinforcement Learning: An Introduction” by Richard Sutton and Andrew Barto
- The definitive RL reference
- Some of these slides were adapted from Pascal Poupart’s slides (CS686 U.Waterloo)
- The TD demo is from Andrej Karpathy